

# Introduction to Abel and queuing system

April 26 2018 - INF9380

Sabry Razick PhD.

Senior Engineer

Research Infrastructure Services Group,  
Department for Research Computing, USIT

# Research Infrastructure Services Group


- Gruppe for forskningsinfrastruktur (FI)
- The FI group provides access to IT resources and high performance computing to researchers at UiO and to NOTUR users
- <http://www.usit.uio.no/om/organisasjon/itf/fi/index.html>
- Part of USIT
- Contact: [hpc-drift@usit.uio.no](mailto:hpc-drift@usit.uio.no)



# Topics

- FI Group
- Abel details
- Getting an account & Logging in
- SLURM Workload Manager
- Running a simple job
- Jobscripts
  - Arrayrun
  - srun
  - MPI
  - GPU

- Nodes - 700+ (704),
- Cores - 10000+ (11,392),  
258 Teraflops
- Total memory - 50 TiB+  
(50.78 TebiBytes)
- Total storage - 400 TiB  
using BeeGFS
- Some nodes with NVIDIA  
Kepler II GPUs (c19-9)
- 96th most powerful in 2012  
, now 444th (June 2015)




**Niels Henrik Abel**

# The Research Computing Services

- Operation of Abel and Colossus computer clusters
- Joint operation of the Fram cluster and NIRD
- User support
- Data storage
- Teaching/Training
- Participation in Nordic and European projects

# Abel Cluster

- Hardware
  - Powerful computers(nodes)
  - Nodes are similar
  - High-speed interconnection
  - Access to a common file system
  - Tightly coupled
- Software
  - Operating system Linux, 64 bit Centos 6.8 (Rocks Cluster Distribution based)
  - Identical mass installations.
  - Queuing system enables timely execution of many concurrent processes
- How it is different from a GRID or CLOUD ?



# Accessing Abel

- If you are working or studying at UiO, you can have an Abel account directly from us.
- If you are Norwegian scientist (or need large resources), you can apply through NOTUR – <http://www.sigma2.no>
- Help and info email: [hpc-drift@usit.uio.no](mailto:hpc-drift@usit.uio.no)




# Connecting to Abel

- Linux
  - Redhat - RHEL
  - Ubuntu
- Windows - using Git BASH, Putty and WinSCP
- Mac OS



# File transfer on command line

- When you log into Abel you are in one of the login nodes login0 - login3.
- Please **DO NOT** execute programs (jobs) directly on the login nodes.
- Jobs are submitted to Abel via the **queuing system**.
- The login nodes are for
  - logging in, copying files, editing, compiling, submitting jobs, checking job status, etc.
- For interactive execution use **qlogin**.



# File transfer

- Unix users can use secure copy or rsync commands
  - For large files, use rsync command
  - Many large files use parallel rsync
    - module load parsync
  - For large number of file use archives
  - E.g. Copy myfile.txt from the current directory on your machine to your home area on Abel:
    - `scp myfile.txt user_name@abel.uio.no:`
- Windows users we recommend WinSCP

# Software on Abel


- Available on Abel:
  - <http://www.uio.no/hpc/abel/help/software>
- Software on Abel are organized in to modules.
- List all software (and version) organized in modules:

```
> module avail
```

- Load software from a module:

```
> module load <module_name>
```

Module load



Blast 2.2.6

Samtools 1.2

Python 2.7.10

# Queue management - SLURM

- Simple Linux Utility for Resource Management (workload manager)
  - Allocates exclusive and/or non-exclusive access to resources (computer nodes) to users for some duration of time
  - Provides a framework for starting, executing, and monitoring work on a set of allocated nodes.
  - Managing a queue of pending work.



# Fair resource allocation

- When you request resources, SLURM on Abel will consider number of things before granting it
  - Does your project has enough CPU hours to pay for this. It will consider total allocated + reserved when doing this.
  - Can/should the cluster provide you with resources (resource combination)
  - Depending on the current load. how long others need to wait if you job starts.

# Laptop Vs queue

- Double click on an icon, give parameters or upload data
  - Wait until execution is finished
  - Inspect results
- Terminal
  - `./myprg input1.txt out_put.file`
  - Inspect results



# Project account

- Needs a project account.
  - UiO students and employees assigned to project **uio** by default.  
But this has a long queue
  - Recommended to apply for own project account from Sigma  
(<https://www.sigma2.no/content/apply-e-infrastructure-resources>)
  - Use the following command to find out the projects that you have access to

> **projects**



# Running a job on the cluster -1



- Login to Abel from you laptop
- Request to occupy some resources from SLURM
- Wait until SLURM grant you the resources
- Execute the job as it was in your laptop

# Interactive use of Abel - qlogin


- Send request for a resource (e.g. 4 cores)
- Work on command line when a node becomes available
- Example - book one node on Abel for your interactive use for 1 hour:

**qlogin --account=your\_project --nodes=1 --exclusive --time=01:00:00**

- Run “**source /cluster/bin/jobsetup**” after receiving allocation
- For more info, see:

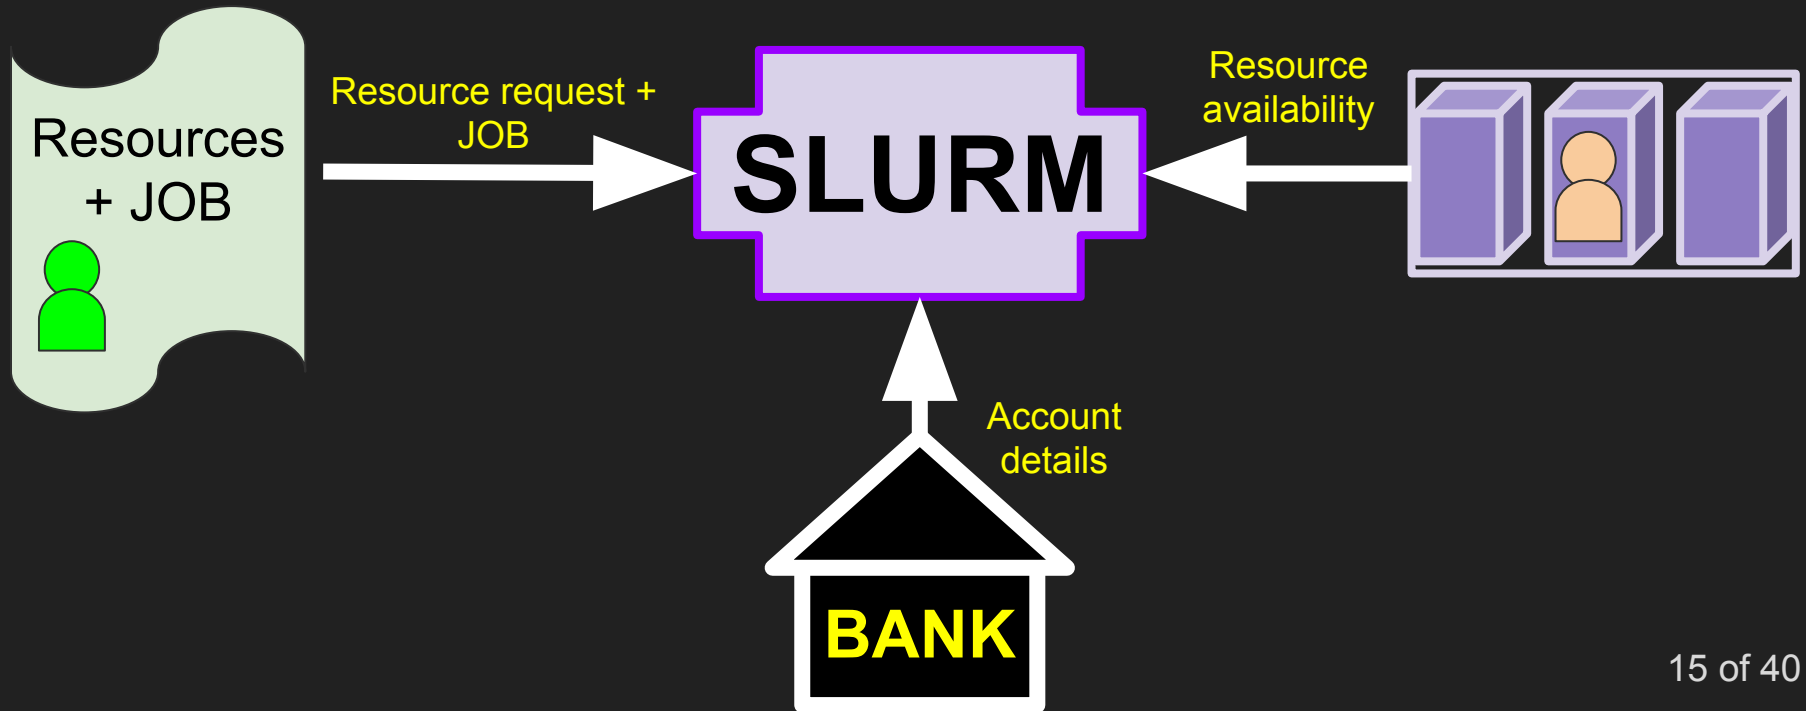
<http://www.uio.no/hpc/abel/help/user-guide/interactive-logins.html>

# Running a job on the cluster - 2

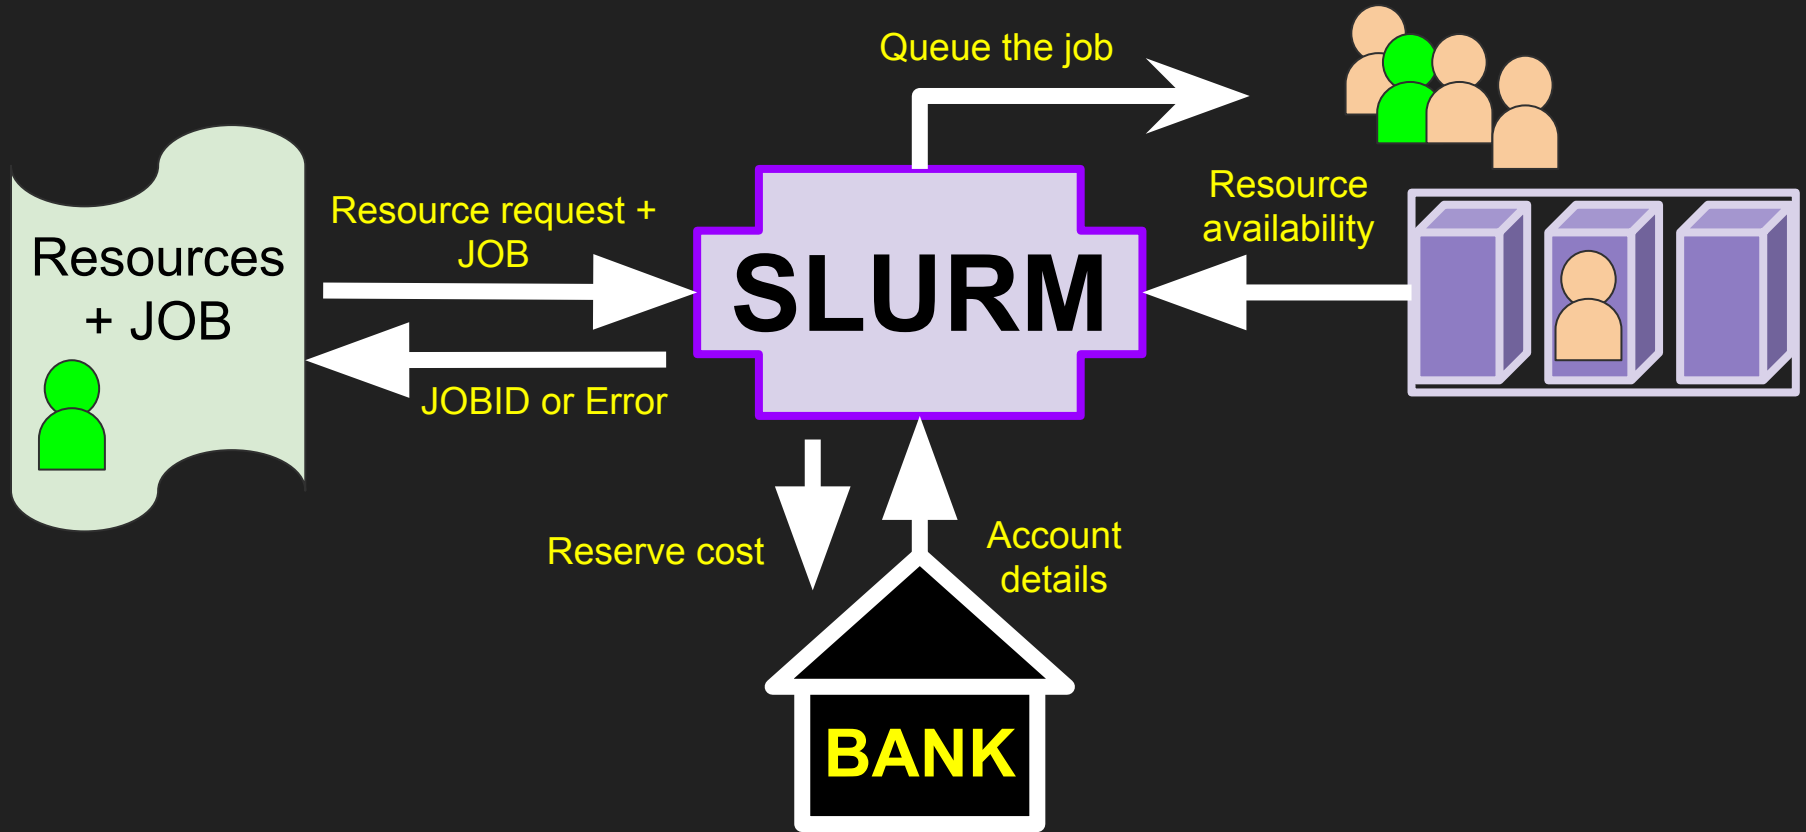


- Login to Abel from you laptop
- Create a job script, with parameters and include the program to run
- Hand it over to the workload manager
- The workload manager will handle the job queue, monitor the progress and let you know the outcome.
- Inspect results

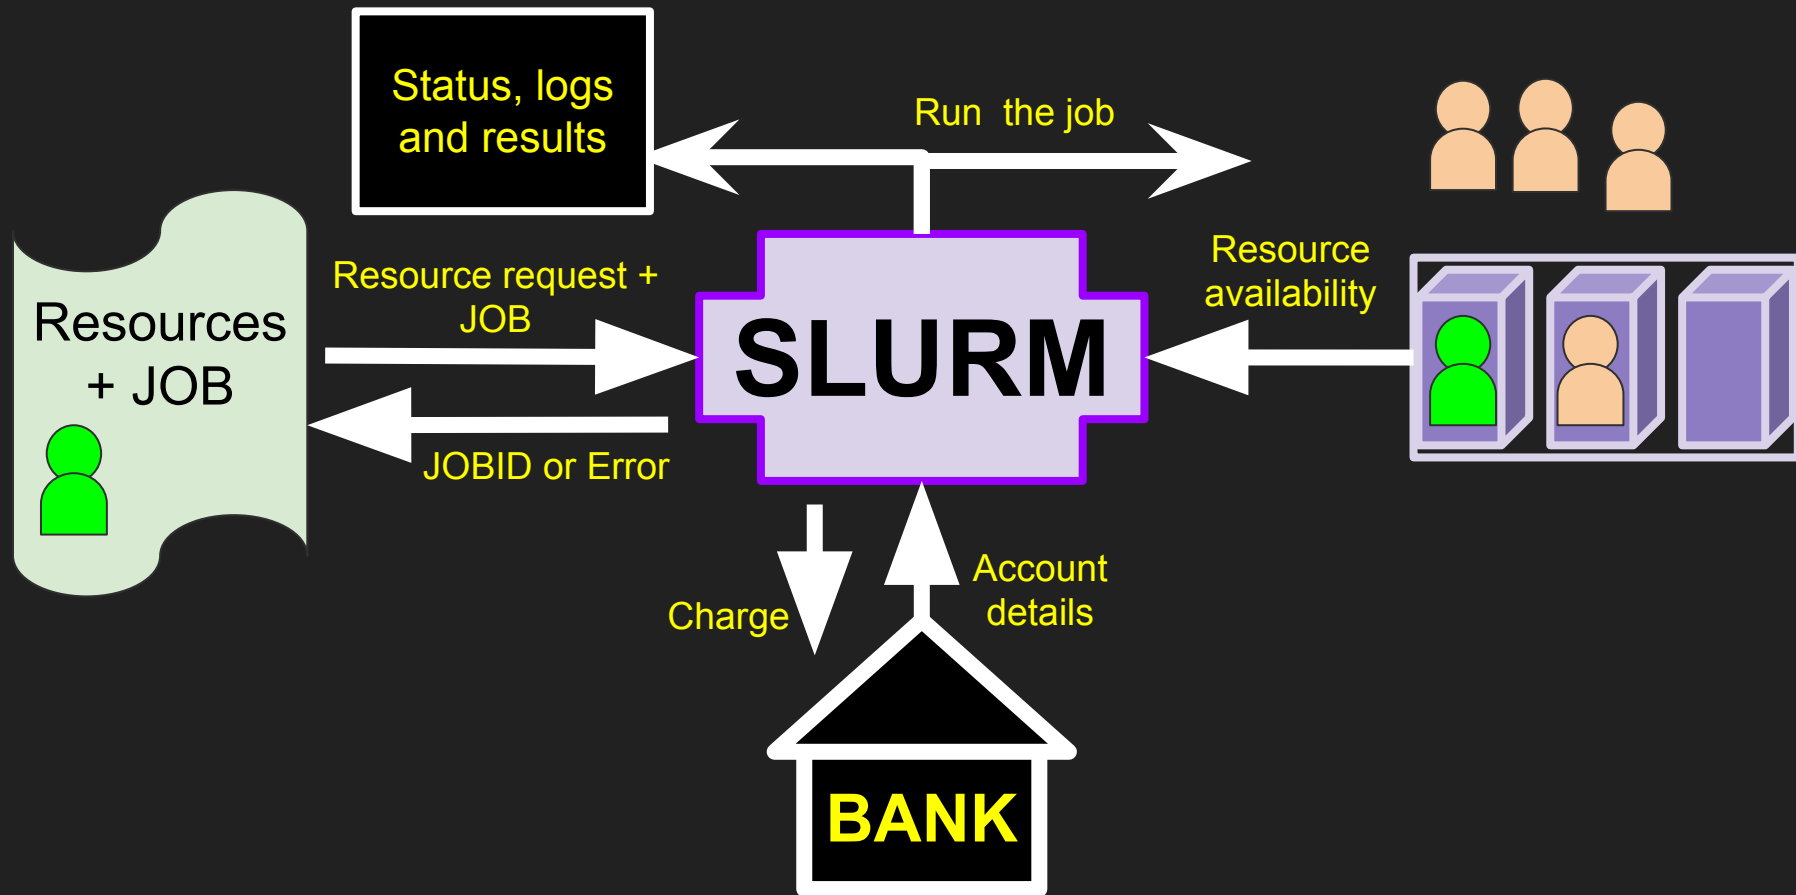
# Job submission to Abel



# Job submission to Abel



# Job submission to Abel



# Computing on Abel

- Get an account and one or more projects (Sigma)
- Load modules or install software
- Submit a job to the queuing system
- Communicate with the queuing system
- Retrieve results (or errors) when the job is done
- Read tutorial: <http://www.uio.no/hpc/abel/help/user-guide>

# Job script

- Job script - shell script including the command that one needs to execute.



## Job script

```
#!/bin/bash
```

```
#SBATCH --job-name=norbis_hello
```

```
#SBATCH --account=ln0002k
```

```
#SBATCH --ntasks=1
```

```
#SBATCH --mem-per-cpu=512M
```

```
#SBATCH --time=00:01:00
```

```
source /cluster/bin/jobsetup
```

```
python hello.py
```

# Job script

```
#!/bin/bash
```

```
#SBATCH --job-name=norbis_hello  
#SBATCH --account=ln0002k  
#SBATCH --ntasks=1  
#SBATCH --mem-per-cpu=512M  
#SBATCH --time=00:01:00
```

Resources

```
source /cluster/bin/jobsetup
```

Setup

```
python hello.py
```


Job

# Understand resources

- If program can not or would not use parallel processing, no point requesting more than one core.
- If you request one tasks you will get exactly one core, from one processor in a single compute node.
  - What does this mean ?



**Supermicro X9DRT compute nodes**



**Samsung DDR3  
1600 MHz**




**Intel® Xeon® Processor  
E5-2670  
Octa core (8)**

Datasheet:


[http://ark.intel.com/products/64595/Intel-Xeon-Processor-E5-2670-20M-Cache-2\\_60-GHz-8\\_00-GTs-Intel-QPI](http://ark.intel.com/products/64595/Intel-Xeon-Processor-E5-2670-20M-Cache-2_60-GHz-8_00-GTs-Intel-QPI)



**Supermicro X9DRT compute nodes**




**Samsung DDR3  
1600 MHz**



**Intel® Xeon® Processor  
E5-2670  
Octa core (8)**

Datasheet:





[http://ark.intel.com/products/64595/Intel-Xeon-Processor-E5-2670-20M-Cache-2\\_60-GHz-8\\_00-GTs-Intel-QPI](http://ark.intel.com/products/64595/Intel-Xeon-Processor-E5-2670-20M-Cache-2_60-GHz-8_00-GTs-Intel-QPI)



- Supermicro X9DRT compute node
- 2 X Intel Xeon E5-2670 (Sandy Bridge) based running at 2.6 GHz,
- 16 physical compute cores total per node.
- Each node have 64 GiBytes of Samsung DDR3 memory


# #SBATCH --ntasks=8

Possible combinations



```
#SBATCH --nodes=1  
#SBATCH --ntasks-per-node=16
```

```
#SBATCH --ntasks=2  
#SBATCH --cpus-per-task=8  
8 X 2 = 16
```





# Time

- Requesting long time will make the queue waiting time longer
- Maximum 1 week (7:00:00:00), hugemem or long partitions 4 weeks

# Memory

- Can request in megabytes or gigabytes
- `#SBATCH --mem-per-cpu=1024M`
- `#SBATCH --mem-per-cpu=1G`
- Request is per CPU (following will reserve 4Gb in total)
  - `#SBATCH --cpus-per-task=4`
  - `#SBATCH -ntask=1`
  - `#SBATCH --mem-per-cpu=1G`
- When requesting a lot of memory you may occupy more cores than requested (following will occupy 16 cores on a regular node)
  - `#SBATCH -ntask=1`
  - `#SBATCH --mem-per-cpu=61G`

# Core hours calculation

```
KMEM=4580.2007628294
```

```
(/cluster/var/accounting/PE_factor)
```

```
PE= NumCPUs
```

```
if(MinMemoryCPU>KMEM){
```

```
    PE=PE*(MinMemoryCPU/KMEM)
```

```
}
```

```
PE_hours = $PE * TimeLimit / 3600
```

# Core hours calculation

#SBATCH --nodes=1

#SBATCH --time=01:00:00

#SBATCH --ntasks-per-node=4

#SBATCH --mem-per-cpu=15G

$(4 \times 1) + 12 = 16$

\*All memory occupied,

KMEM=4580.2007628294

PE= 4

#(15 \* 1024)>KMEM so

$PE = 4 * ((15 * 1024) / KMEM) = 13.41$

$PE\_hours = 13.41 * (1 * 60 * 60) / 3600 = 13.41$

\*\*Use the command **cost** to check account balance

**\$ cost**

# Sample job script

```
hello.slurm
```

```
#!/bin/bash
```

```
#SBATCH --job-name=RCS1115_hello
```

```
#SBATCH --account=xxx
```

```
#SBATCH --ntasks=1
```

```
#SBATCH --mem-per-cpu=256M
```

```
#SBATCH --time=00:05:05
```

```
source /cluster/bin/jobsetup
```

```
set -o errexit
```

```
sleep 5m
```

```
echo "Hello from "
```

```
hostname
```

## Join the queue

```
> sbatch hello.slurm
```

```
> scontrol show job <JOBID>
```

```
> squeue -u <username>
```

```
> sacct -j <JOBID>
```

# Investigate

```
>scontrol show job 12989353
  JobId=12989353 Name=RCS1115_hello
  UserId=sabryr(243460) GroupId=users(100)
  Priority=22501 Nice=0 Account=staff      QOS=staff
  JobState=COMPLETED Reason=None
...
  RunTime=00:00:02 TimeLimit=00:01:00 Ti
Command=../RCS_tutorial/hello.slurm
WorkDir=../RCS_tutorial
StdErr=../RCS_tutorial/slurm-12989552.out
StdOut=../RCS_tutorial/slurm-12989552.out
```

## Some usefull commands

- **scancel <JOBID>** - Cancel a job before it ends
- **dusage** - find out your disk usage
- **squeue** - list all queued jobs and find out the
- **squeue -t R | more** -position of your job
- **cost** - account balance
- **sinfo** - node info
- **sacct** - details of jobs that have ended



# Environment variables

- SLURM\_JOBID – job-id of the job
- SCRATCH – name of job-specific scratch-area
- SLURM\_NPROCS – total number of cpus requested
- SLURM\_CPUS\_ON\_NODE – number of cpus allocated on node
- SUBMITDIR – directory where sbatch were issued
- SCRATCH - scratch directory location (on /work)
- TASK\_ID – task number (for arrayrun-jobs)

# Check Variables

```
checkvariables.slurm
#!/bin/bash

#SBATCH --job-name=RCS1115_hello
#SBATCH --account=staff
#SBATCH --time=00:01:05
#SBATCH --mem-per-cpu=512M
#SBATCH --ntasks=1

source /cluster/bin/jobsetup

echo "SLURM_JOBID=" $SLURM_JOBID
echo "SCRATCH=" $SCRATCH
echo "SLURM_NPROCS=" $SLURM_NPROCS
echo "SLURM_CPUS_ON_NODE=" $SLURM_CPUS_ON_NODE
echo "SUBMITDIR=" $SUBMITDIR
echo "TASK_ID=" $TASK_ID
```

Parallel jobs

# Arrayrun

- Every arrayrun needs one “worker” script and one “submit” script.
- Worker script - similar to the jobscript we have used
- The submit script executes the arrayrun command
- TASK\_ID variable is used to distinguish each run

# Arrayrun - submit script

```
array_submit.slurm
#!/bin/bash

#SBATCH --job-name=RCS1115_hello
#SBATCH --account=staff
#SBATCH --time=00:02:00
#SBATCH --ntasks=1
#SBATCH --mem-per-cpu=256M

source /cluster/bin/jobsetup
set -o errexit

module purge
module load Python/3.5.2-foss-2016b

arrayrun 1-3 array_worker.slurm
```

# Arrayrun - submit script

```
array_worker.slurm
#!/bin/bash

#SBATCH --job-name=RCS1115_hello
#SBATCH --account=staff
#SBATCH --time=00:02:00
#SBATCH --ntasks=1
#SBATCH --mem-per-cpu=256M

source /cluster/bin/jobsetup
set -o errexit

module purge
module load Python/3.5.2-foss-2016b

arrayrun 1-3 array_worker.slurm
```

# Investigate job

```
> ssh to compute node and use top
```

```
> ssh c16-17 "top -n1 -b | grep <USER_NAME>"
```

```
> ssh c60-1 "lscpu | grep Socket"
```

# srun

```
srun_test.slurm
#!/bin/bash

#SBATCH --job-name=RCS1115_hello
#SBATCH --account=staff
#SBATCH --time=00:01:00
#SBATCH --ntasks=3
#SBATCH --mem-per-cpu=256M

FILES=(/path/to/data/*)

source /cluster/bin/jobsetup
set -o errexit

echo "number of instances "$SLURM_NTASKS
srun aprogram.sh
```



# MPI

- Message Passing Interface
- Especially when running on more than one node
- MPI is a language-independent communications protocol used for programming parallel computers.
- jobs specifying more than one node automatically get
  - `#SBATCH --constraint=ib`
- We OpenMPI and Intel® MPI on Abel
  - `module load openmpi.gnu/1.10.2`
  - `module load openmpi.intel/1.10.2`
  - `module load intelmpi.intel/5.0.2`

# MPI

```
mpi_test.slurm
#!/bin/bash

#SBATCH --job-name=RCS1115_hello
#SBATCH --account=staff
#SBATCH --time=00:01:00
#SBATCH --ntasks=3
#SBATCH --mem-per-cpu=256M

source /cluster/bin/jobsetup
set -o errexit

module purge
module load Python/3.5.2-foss-2016b

mpirun python mpi_test.py
```

GPU

# GPU job

```
gpu.slurm
#!/bin/bash

#SBATCH --job-name=GPU_test
#SBATCH --account=ln0002K
#SBATCH --ntasks=1
#SBATCH --time=01:00
#SBATCH --mem-per-cpu=1G
## Ask for 1 GPU
#SBATCH --partition=accel --gres=gpu:1

## Set up job environment:
source /cluster/bin/jobsetup
set -o errexit # exit on errors

##Print number of available GPUs and some details about them
echo $CUDA_VISIBLE_DEVICES
nvidia-smi --query-gpu=name,memory.total,memory.free,memory.used --format=csv
```

Visualization

# Visualization example

- `ssh -X abel.uio.no`
- `qlogin -A staff --ntasks=2 --mem-per-cpu=2G --time=10:00 --job-name=R_test`
- `module load R/3.4.4`
- `R`
- `> x <- c(1:5); y <- x`
- `> plot(x, y, type="p")`

I/O

# \$SCRATCH

- If you are accessing a file multiple time during a job
  - Use scratch directory
  - Use /work/users/<USERNAME>
- Choose where you start the job from
- On Abel jobs accessing files and/or writing out large outputs will run faster if /work is used compared to running from \$HOME
- When running jobs from /work, there is no need use \$SCRATCH
- Remember to use to chkfile copy back results, otherwise the results will be deleted at the end of the job.



# \$SCRATCH

```
scratch.slurm
#!/bin/bash

#SBATCH --job-name=usescratch
...

source /cluster/bin/jobsetup
set -o errexit

echo $SUBMITDIR
echo $SCRATCH
cp $SUBMITDIR/input* $SCRATCH
chkfile output.txt
cat input* > output.txt
```

# \$LOCALTMP

```
Localtmp.slurm & fromhome.slurm
```

```
#!/bin/bash
```

```
....
```

```
source /cluster/bin/jobsetup
```

```
set -o errexit
```

```
echo $SUBMITDIR
```

```
echo $LOCALTMP
```

```
df -h $LOCALTMP
```

```
cp manyfiles.sh $LOCALTMP
```

```
cd $LOCALTMP
```

```
pwd
```

```
./manyfiles.sh
```

```
rsync all.txt $SUBMITDIR/
```

# I/O - Files

- When handling very large number of files try to use
  - Archives - just in time extract or write directly to archive
  - Cleanup unwanted files
  - Copy back only output files needed when using `$SCRATCH` , e.g, do not copy back input data
- Get advice on using `/tmp` directory on compute nodes when thousands of files.

# I/O - Files

- Create an archive (no compression)
  - `tar -cvf <ARCHIVE_NM> <FILES>`
  - `tar -cvf file.tar *.txt`
- List content
  - `tar -tvf <ARCHIVE_NM>`
- Append a file
  - `tar --append --file <ARCHIVE_NM> <NEW_FILES>`
  - `tar --append --file files.tar.gz 1.txt`
- Extract all
  - `tar -xvzf <ARCHIVE_NM>`
- Extract one
  - `tar -xvf <ARCHIVE_NM> <FILES>`

# Thank you.



<http://www.uio.no/english/services/it/research/hpc/abel/>



[hpc-drift@usit.uio.no](mailto:hpc-drift@usit.uio.no)

# sbatch - memory

- #SBATCH --mem-per-cpu=Size
  - Memory required per allocated core (format: 2G or 2048M)
  - How much memory should one specify? The maximum usage of RAM by your program (plus some). Exaggerated values might, delay the job start.
- #SBATCH --partition=hugemem
  - If you need more than 61GB of RAM on a single node (up to 1 TiB).

# sbatch - time

- `#SBATCH --time=hh:mm:ss`
  - Wall clock time limit on the job
  - Some prior testing is necessary. One might, for example, test on smaller data sets and extrapolate. As with the memory, unnecessarily large values might delay the job start.
  - This costs you (from allocated operation resources)
  - Until a job is finished this will be reserved.
- `#SBATCH --begin=hh:mm:ss`
  - Start the job at a given time (or later)
- `#SBATCH --partition=long`
  - Maximum time for a job is 1 week (168 hours). If more needed, use

# sbatch – CPUs and nodes

- Does your program support more than one CPU?
- If so, do they have to be on a single node?
- How many CPUs will the program run efficiently on?
- #SBATCH --nodes=Nodes
  - Number of nodes to allocate
- #SBATCH --ntasks-per-node=Cores
  - Number of cores to allocate within each allocated node
- #SBATCH --ntasks=Cores
  - Number of cores to allocate
- #SBATCH --cpus-per-task=Cores
  - Threads on one node



# sbatch - interconnect

- #SBATCH --constraint=*ib*
  - Run job on nodes with infiniband
  - Gigabit Ethernet on all nodes
  - All nodes on Abel are equipped with InfiniBand (56 Gbits/s)
  - Select if you run MPI jobs

# sbatch - constraint

- #SBATCH --constraint=*feature*
  - Run job on nodes with a certain feature - *ib, rackN*.
- #SBATCH --constraint=*ib&rack21*
  - If you need more than one constraint
  - *in case of multiple specifications, the later overrides the earlier*

# sbatch - files

- #SBATCH --output=file
  - Send 'stdout' (and stderr) to the specified file (instead of slurm-xxx.out)
- #SBATCH --error=file
  - Send 'stderr' to the specified file
- #SBATCH --input=file
  - Read 'stdin' from the specified file

# sbatch – low priority

- #SBATCH --qos=lowpri
  - Run a job in the lowpri queue
  - Even if all of your project's cpus are busy, you may utilize other cpus
  - Such a job may be terminated and put back into the queue at any time.
  - If possible, your job should ensure its state is saved regularly, and should be prepared to pick up on where it left off.
  - Note: Notur projects cannot access lowpri.