

# Parallel programming in R

Bjørn-Helge Mevik

Research Infrastructure Group, USIT, UiO

NORBIS Course Week, april 2016

Introduction

Simple example

Practical use

The end...

# Background

- ▶ R is single-threaded
- ▶ There are several packages for parallel computation in R, some of which have existed a long time, e.g. `Rmpi`, `nws`, `snow`, `sprint`, `foreach`, `multicore`
- ▶ As of 2.14.0, R ships with a package `parallel`
- ▶ R can also be compiled against multi-threaded linear algebra libraries (BLAS, LAPACK) which can speed up some calculations

Today's focus is the `parallel` package.

# Overview of `parallel`

- ▶ Introduced in 2.14.0
- ▶ Based on packages `multicore` and `snow` (slightly modified)
- ▶ Includes a parallel random number generator (RNG); important for simulations (see `?nextRNGStream`)
- ▶ Particularly suitable for 'single program, multiple data' (SPMD) problems
- ▶ Main interface is parallel versions of `lapply` and similar
- ▶ Can use the CPUs/cores on a single machine (`multicore`), or several machines, using MPI (`snow`)
- ▶ MPI support depends on the `Rmpi` package (installed on Abel)

## Simple example: serial

- ▶ `parallel` provides substitutes for `lapply`, etc.
- ▶ 'Silly' example for illustration: calculate  $(1:100)^2$

Serial version:

```
## The worker function to do the calculation:
```

```
workerFunc <- function(n) { return(n^2) }
```

```
## The values to apply the calculation to:
```

```
values <- 1:100
```

```
## Serial calculation:
```

```
res <- lapply(values, workerFunc)
```

```
print(unlist(res))
```

## Simple example: `mclapply`

- ▶ Performs the calculations in parallel on the local machine
- ▶ (+) Very easy to use; no set-up
- ▶ (+) Low overhead
- ▶ (-) Can only use the cores of *one* machine
- ▶ (-) Uses fork, so it will not work on MS Windows

```
workerFunc <- function(n) { return(n^2) }  
values <- 1:100
```

```
library(parallel)
```

```
## Number of workers (R processes) to use:  
numWorkers <- 8
```

```
## Parallel calculation (mclapply):  
res <- mclapply(values, workerFunc, mc.cores = numWorkers)
```

```
print(unlist(res))
```

## Simple example: `parLapply`

- ▶ Performs the calculations in parallel, possibly on several nodes
- ▶ Can use several types of communications, including `PSOCK` and `MPI`
- ▶ `PSOCK`:
  - ▶ (+) Can be used interactively
  - ▶ (-) Not good for running on several nodes
  - ▶ (+) Portable; works 'everywhere'
  - ▶ ==> Good for testing
- ▶ `MPI`:
  - ▶ (-) Needs the `Rmpi` package (installed on Abel)
  - ▶ (-) Cannot be used interactively
  - ▶ (+) Good for running on several nodes
  - ▶ (+) Works everywhere where `Rmpi` does
  - ▶ ==> Good for production

## Simple example: `parLapply` (PSOCK)

```
workerFunc <- function(n) { return(n^2) }
values <- 1:100

library(parallel)

## Number of workers (R processes) to use:
numWorkers <- 8

## Set up the 'cluster'
cl <- makeCluster(numWorkers, type = "PSOCK")

## Parallel calculation (parLapply):
res <- parLapply(cl, values, workerFunc)

## Shut down cluster
stopCluster(cl)

print(unlist(res))
```



## Simple example: `parLapply` (MPI)

`simple_mpi.R`:

```
workerFunc <- function(n) { return(n^2) }  
values <- 1:100  
library(parallel)  
numWorkers <- 8  
cl <- makeCluster(numWorkers, type = "MPI")  
res <- parLapply(cl, values, workerFunc)  
stopCluster(cl)  
Rmpi::mpi.finalize() # or Rmpi::mpi.quit(), which quits R as well  
print(unlist(res))
```

Running:

```
mpirun -n 1 R --slave -f simple_mpi.R
```

Note: Use R  $\geq$  2.15.2 for [MPI](#), due to a bug in earlier versions of [parallel](#).

# Preparation for calculations

- ▶ Write your calculations as a function that can be called with `lapply`
- ▶ Test interactively with `lapply` serially, and `mclapply` or `parLapply` (`PSOCK`) in parallel
- ▶ Deploy with `mclapply` on single node or `parLapply` (`MPI`) on one or more nodes
- ▶ For `parLapply`, the worker processes must be prepared with any loaded packages with `clusterEvalQ` or `clusterCall`.
- ▶ For `parLapply`, large data sets can be exported to workers with `clusterExport`.

## Extended example

(Notes to self:)

- ▶ Submit jobs
- ▶ Go through scripts
- ▶ Look at results

# Efficiency

- ▶ The time spent in each invocation of the worker function should be long enough
- ▶ If the time spent in each invocation of the worker function vary very much, try the load balancing versions of the functions
- ▶ Avoid copying large things back and forth:
  - ▶ Export large datasets up front with `clusterExport` (for `parLapply`)
  - ▶ Iterate over indices or similar small things, not large data sets
  - ▶ Let the worker function return as little as possible
- ▶ Reduce waiting time in queue by not asking for whole nodes. If possible, use `--ntask` instead of `--ntasks-per-node` + `--nodes` (this means using MPI).

## Other topics

There are several things we haven't touched in this lecture:

- ▶ Parallel random number generation
- ▶ Alternatives to `*apply` (e.g. `mcpipeline` + `mccollect`)
- ▶ Lower level functions
- ▶ Using multi-threaded libraries
- ▶ Other packages and techniques

Resources:

- ▶ The documentatin for `parallel`: `help(parallel)`
- ▶ The book *Parallel R*, McCallum & Weston, O'Reilly
- ▶ The HPC Task view on CRAN:  
<http://cran.r-project.org/web/views/HighPerformanceComputing.html>