

Parallel programming in R

Bjørn-Helge Mevik

Dept. for Research Infrastructure Services, USIT, UiO

INF9380, April 2020

Introduction

Simple example

Practical use

The end...

Background

- ▶ R is (more-or-less) single-threaded
- ▶ There are several contributed packages for parallel computation in R, some of which have existed a long time, e.g. `Rmpi`, `nws`, `snow`, `sprint`, `foreach`, `multicore`
- ▶ R itself ships with a package `parallel`
- ▶ R can also be compiled against multi-threaded linear algebra libraries (BLAS, LAPACK) which can speed up some calculations

Today's focus is the `parallel` package.

To Parallelise or not to Parallelise?

- ▶ Parallelisation is all about increasing speed
- ▶ It can be tricky to get good speed increase by parallelisation
 - ▶ Overhead
 - ▶ Algorithmic constraints
- ▶ Start with optimising your code serially and linking with optimised libraries
- ▶ If you still need more speed: parallelise

What Level to Parallelise?

General rules of thumb:

- ▶ The higher level, the better
- ▶ The more independent chunks of calculation, the better
- ▶ Library level can be very easy to implement (but might not give the same speed increase)
- ▶ One level is usually better than two levels

How Wide to Parallelise?

- ▶ If you have one calculation to perform, and a number of CPUs that would otherwise be idle: Increase number of cpus until walltime starts to increase.
- ▶ In most other situations: Increase number of cpus until cputime (walltime \times nCPUs) increases too much.
- ▶ If you are developing an R-package: Let the users choose.

Overview of `parallel`

- ▶ Uses *processes*, not *threads*.
- ▶ Based on packages `multicore` and `snow` (slightly modified)
- ▶ Includes a parallel random number generator (RNG); important for simulations (see `?nextRNGStream`)
- ▶ Particularly suitable for 'single program, multiple data' (SPMD) problems
- ▶ Main interface is parallel versions of `lapply` and similar
- ▶ Can use the CPUs/cores on a single machine (`multicore/snow`), or several machines, using MPI (`snow`)
- ▶ MPI support depends on the `Rmpi` package (installed on Saga)

Simple example

- ▶ Silly example for illustration: calculate $(1:100)^2$

Simple example: serial

- ▶ `parallel` provides substitutes for `lapply`, etc.

Serial version:

```
## The worker function to do the calculation:
```

```
workerFunc <- function(n) { return(n^2) }
```

```
## The values to apply the calculation to:
```

```
values <- 1:100
```

```
## Serial calculation:
```

```
res <- lapply(values, workerFunc)
```

```
print(unlist(res))
```

Simple example: `mclapply`

- ▶ Performs the calculations in parallel on the local machine
- ▶ (+) Very easy to use; no set-up
- ▶ (+) Can be used interactively
- ▶ (+) Less overhead
- ▶ (-) Can only use the cores of *one* machine
- ▶ (-) Uses fork, so it will not work on MS Windows
- ▶ (-) Can leave orphaned R processes
- ▶ => Good for within-one machine, but perhaps not for repeated calculations

Simple example: `mclapply`

```
workerFunc <- function(n) { return(n^2) }  
values <- 1:100  
  
library(parallel)  
  
## Number of workers (R processes) to use:  
numWorkers <- 8  
  
## Parallel calculation (mclapply):  
res <- mclapply(values, workerFunc, mc.cores = numWorkers)  
  
print(unlist(res))
```

Simple example: `parLapply`

- ▶ Performs the calculations in parallel, possibly on several machines
- ▶ Can use several types of communications, including `PSOCK` and `MPI`
- ▶ `PSOCK`:
 - ▶ (+) Can be used interactively
 - ▶ (+) Portable; works 'everywhere'
 - ▶ (-) Not good for running on several machines
 - ▶ ==> Good for within-one-machine and tests
- ▶ `MPI`:
 - ▶ (-) Needs the `Rmpi` package (installed on Saga)
 - ▶ (-) Cannot be used interactively
 - ▶ (+) Good for running on several machines
 - ▶ (+) Works everywhere where `Rmpi` does
 - ▶ ==> Good for multi-machine

Simple example: `parLapply` (PSOCK)

```
workerFunc <- function(n) { return(n^2) }
values <- 1:100

library(parallel)

## Number of workers (R processes) to use:
numWorkers <- 8

## Set up the 'cluster'
cl <- makeCluster(numWorkers, type = "PSOCK")

## Parallel calculation (parLapply):
res <- parLapply(cl, values, workerFunc)

## Shut down cluster
stopCluster(cl)

print(unlist(res))
```

Simple example: `parLapply` (MPI)

`simple_mpi.R`:

```
workerFunc <- function(n) { return(n^2) }  
values <- 1:100  
library(parallel)  
numWorkers <- 8  
cl <- makeCluster(numWorkers, type = "MPI")  
res <- parLapply(cl, values, workerFunc)  
stopCluster(cl)  
Rmpi::mpi.finalize() # or Rmpi::mpi.quit(), which quits R as well  
print(unlist(res))
```

Running:

```
mpirun -n 1 R --slave -f simple_mpi.R
```

Preparation for calculations

- ▶ Write your calculations as a function that can be called with `lapply`
- ▶ Test interactively with `lapply` serially, and `mclapply` or `parLapply` (`PSOCK`) in parallel
- ▶ Deploy with `mclapply` on single machine or `parLapply` on one or more machines
- ▶ For `parLapply`, the worker processes must be prepared with any loaded packages with `clusterEvalQ` or `clusterCall`.
- ▶ For `parLapply`, large data sets can be exported to workers with `clusterExport`.

Extended example: Cross-Validation

Cross-validation of regression model.

(Notes to self:)

- ▶ Submit jobs
- ▶ Go through scripts
- ▶ Look at results

Efficiency

- ▶ Don't use more processes (or threads) than you have *physical* CPU cores!
- ▶ Avoid copying large things back and forth:
 - ▶ Export large datasets up front with `clusterExport` (for `parLapply`)
 - ▶ Iterate over indices or similar small things, not large data sets
 - ▶ Let the worker function return as little as possible
- ▶ The time spent in each invocation of the worker function should be long enough
- ▶ If the time spent in each invocation of the worker function vary very much, try the load balancing versions of the functions

Illustration: Don't use logical CPU cores

- ▶ Machine with 4 physical cores, 2 logical cpus per core

- ▶ Worker functions:

```
sleeper <- function(n) {
  Sys.sleep(10)
  return(42)
}
worker <- function(n) {
  i <- 1
  while(i < 2.5e7) { tmp <- sqrt(42); i <- i + 1 }
  return(42)
}
```

- ▶ Timings:

```
> system.time(sleeper(1))
  user  system elapsed
0.000   0.000  10.006
> system.time(worker(1))
  user  system elapsed
9.948   0.000   9.949
```

Illustration: Don't use logical CPU cores

► No real computation:

```
> system.time(mclapply(1:2, sleeper, mc.cores = 2))
  user  system elapsed
0.008   0.000  10.013
> system.time(mclapply(1:4, sleeper, mc.cores = 4))
  user  system elapsed
0.004   0.004  10.014
> system.time(mclapply(1:8, sleeper, mc.cores = 8))
  user  system elapsed
0.000   0.008  10.015
> system.time(mclapply(1:16, sleeper, mc.cores = 16))
  user  system elapsed
0.004   0.008  10.015
```

Illustration: Don't use logical CPU cores

► Actual computation:

```
> system.time(mclapply(1:2, worker, mc.cores = 2))
  user  system elapsed 
10.100   0.012  10.122 
> system.time(mclapply(1:4, worker, mc.cores = 4))
  user  system elapsed 
42.048   0.060  10.774 
> system.time(mclapply(1:8, worker, mc.cores = 8))
  user  system elapsed 
160.056   0.176  21.889
```

Illustration: Return as little as possible

▶ Worker functions:

```
small <- function(n) {  
  Sys.sleep(0.01)  
  return(42)  
}  
big <- function(n) {  
  Sys.sleep(0.01)  
  return(1:100000)  
}
```

▶ Timings:

```
> system.time(small(1))  
   user  system elapsed  
  0.00   0.00   0.01  
> system.time(big(1))  
   user  system elapsed  
  0.00   0.00   0.01
```

Illustration: Return as little as possible

► Returning a single number:

```
> system.time(mclapply(1:4000, small, mc.cores = 2))
  user  system elapsed 
0.000   0.004  20.279 

> system.time(mclapply(1:4000, small, mc.cores = 4))
  user  system elapsed 
0.112   0.052  10.135 

> system.time(mclapply(1:4000, small, mc.cores = 8))
  user  system elapsed 
0.152   0.124   5.072 

> system.time(mclapply(1:4000, small, mc.cores = 16))
  user  system elapsed 
0.104   0.040   2.542
```

Illustration: Return as little as possible

► Returing 1×10^5 numbers:

```
> system.time(mclapply(1:4000, big, mc.cores = 2))  
   user  system elapsed  
0.244   0.972  22.898
```

```
> system.time(mclapply(1:4000, big, mc.cores = 4))  
   user  system elapsed  
1.792   2.952  12.034
```

```
> system.time(mclapply(1:4000, big, mc.cores = 8))  
   user  system elapsed  
1.048   2.664   6.922
```

```
> system.time(mclapply(1:4000, big, mc.cores = 16))  
   user  system elapsed  
1.756   3.672   4.296
```

Other topics

There are several things we haven't touched in this lecture:

- ▶ Parallel random number generation
- ▶ Alternatives to `*apply` (e.g. `mcparrallel` + `mccollect`)
- ▶ Lower level functions
- ▶ Using multi-threaded libraries
- ▶ Other packages and techniques

Resources:

- ▶ The documentatin for `parallel`: `library(help = parallel)`, `vignette("parallel")`
- ▶ The book *Parallel R*, McCallum & Weston, O'Reilly
- ▶ The HPC Task view on CRAN:
<http://cran.r-project.org/web/views/HighPerformanceComputing.html>